



МИНОБРНАУКИ РОССИИ
Федеральное государственное бюджетное образовательное учреждение
высшего образования
«МИРЭА – Российский технологический университет»
РТУ МИРЭА

ИКБ направление «Киберразведка и противодействие угрозам с применением технологий искусственного интеллекта» 10.04.01

Кафедра КБ-4 «Интеллектуальные системы информационной безопасности»

Лабораторная работа №1
по дисциплине

«Анализ защищенности систем искусственного интеллекта»

Группа:
ББМО-01-22
Выполнил:
Воронов А.И.

Проверил:
Спирин А.А.

Москва 2023

Выполняем импорт необходимых библиотек.

```
# импортируем библиотеки
import numpy as np
import matplotlib.pyplot as plt
import torch
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim
from torchvision import transforms, datasets
```

Задаем нормализующие преобразования, загружаем набор MNIST и разбиваем данные на подвыборки.

```
# задаем нормализующие преобразования, загружаем набор MNIST и разбиваем данные на подвыборки
transform = transforms.Compose([transforms.ToTensor(),
                                transforms.Normalize((0.0,), (1.0,))])
dataset = datasets.MNIST(root = './data', train=True, transform = transform, download=True)
train_set, val_set = torch.utils.data.random_split(dataset, [50000, 10000])
test_set = datasets.MNIST(root = './data', train=False, transform = transform, download=True)
train_loader = torch.utils.data.DataLoader(train_set, batch_size=1, shuffle=True)
val_loader = torch.utils.data.DataLoader(val_set, batch_size=1, shuffle=True)
test_loader = torch.utils.data.DataLoader(test_set, batch_size=1, shuffle=True)
print("Training data:", len(train_loader), "Validation data:", len(val_loader), "Test data:", len(test_loader))

Downloading http://yann.lecun.com/exdb/mnist/train-images-idx3-ubyte.gz
Downloading http://yann.lecun.com/exdb/mnist/train-images-idx3-ubyte.gz to ./data/MNIST/raw/train-images-idx3-ubyte.gz
100%|██████████| 9912422/9912422 [00:00<00:00, 78744263.55it/s]
Extracting ./data/MNIST/raw/train-images-idx3-ubyte.gz to ./data/MNIST/raw

Downloading http://yann.lecun.com/exdb/mnist/train-labels-idx1-ubyte.gz
Downloading http://yann.lecun.com/exdb/mnist/train-labels-idx1-ubyte.gz to ./data/MNIST/raw/train-labels-idx1-ubyte.gz
100%|██████████| 28881/28881 [00:00<00:00, 92329034.93it/s]
Extracting ./data/MNIST/raw/train-labels-idx1-ubyte.gz to ./data/MNIST/raw

Downloading http://yann.lecun.com/exdb/mnist/t10k-images-idx3-ubyte.gz
Downloading http://yann.lecun.com/exdb/mnist/t10k-images-idx3-ubyte.gz to ./data/MNIST/raw/t10k-images-idx3-ubyte.gz
100%|██████████| 1648877/1648877 [00:00<00:00, 23596879.39it/s]
Extracting ./data/MNIST/raw/t10k-images-idx3-ubyte.gz to ./data/MNIST/raw

Downloading http://yann.lecun.com/exdb/mnist/t10k-labels-idx1-ubyte.gz
Downloading http://yann.lecun.com/exdb/mnist/t10k-labels-idx1-ubyte.gz to ./data/MNIST/raw/t10k-labels-idx1-ubyte.gz
100%|██████████| 4542/4542 [00:00<00:00, 9835069.06it/s]
Extracting ./data/MNIST/raw/t10k-labels-idx1-ubyte.gz to ./data/MNIST/raw

Training data: 50000 Validation data: 10000 Test data: 10000
```

Задаем использование GPU для вычислений, если это возможно.

```
[ ] # используем GPU если возможно
use_cuda=True
device = torch.device("cuda" if (use_cuda and torch.cuda.is_available()) else "cpu")
```

На основе torch создаем класс нейронной сети

```

# создаем класс НС на основе torch
class Net(nn.Module):
    def __init__(self):
        super(Net, self).__init__()
        self.conv1 = nn.Conv2d(1, 32, 3, 1)
        self.conv2 = nn.Conv2d(32, 64, 3, 1)
        self.dropout1 = nn.Dropout2d(0.25)
        self.dropout2 = nn.Dropout2d(0.5)
        self.fc1 = nn.Linear(9216, 128)
        self.fc2 = nn.Linear(128, 10)

    def forward(self, x):
        x = self.conv1(x)
        x = F.relu(x)
        x = self.conv2(x)
        x = F.relu(x)
        x = F.max_pool2d(x, 2)
        x = self.dropout1(x)
        x = torch.flatten(x, 1)
        x = self.fc1(x)
        x = F.relu(x)
        x = self.dropout2(x)
        x = self.fc2(x)
        output = F.log_softmax(x, dim=1)
        return output

```

Проверяем работоспособность созданного класса и создаем оптимизатор, функцию потерь и трейнер сети.

```

# проверяем работоспособность созданного класса
model = Net().to(device)

# создаем оптимизатор, функцию потерь и трейнер сети
optimizer = optim.Adam(model.parameters(), lr=0.0001, betas=(0.9, 0.999))
criterion = nn.NLLLoss()
scheduler = optim.lr_scheduler.ReduceLROnPlateau(optimizer, mode='min', factor=0.1, patience=3)

```

Определяем функцию обучения сети.

```

# определяем функцию обучения сети
def fit(model, device, train_loader, val_loader, epochs):
    data_loader = {'train': train_loader, 'val': val_loader}
    print("Fitting the model...")
    train_loss, val_loss = [], []
    for epoch in range(epochs):
        loss_per_epoch, val_loss_per_epoch = 0, 0
        for phase in ('train', 'val'):
            for i, data in enumerate(data_loader[phase]):
                input, label = data[0].to(device), data[1].to(device)
                output = model(input)
                # calculating loss on the output
                loss = criterion(output, label)
                if phase == 'train':
                    optimizer.zero_grad()
                    # grad calc w.r.t Loss func
                    loss.backward()
                    # update weights
                    optimizer.step()
                    loss_per_epoch += loss.item()
                else:
                    val_loss_per_epoch += loss.item()
            scheduler.step(val_loss_per_epoch / len(val_loader))
        print("Epoch: {} Loss: {} Val_Loss: {}".format(epoch+1, loss_per_epoch / len(train_loader), val_loss_per_epoch / len(val_loader)))
        train_loss.append(loss_per_epoch / len(train_loader))
        val_loss.append(val_loss_per_epoch / len(val_loader))
    return train_loss, val_loss

```

Выполним обучение модели.

```

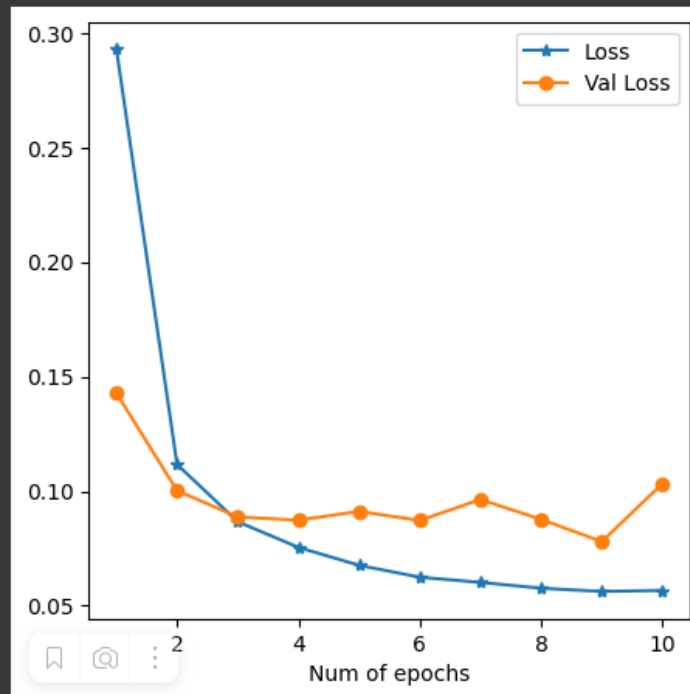
# обучаем модель
loss, val_loss = fit(model, device, train_loader, val_loader, 10)

Fitting the model...
/usr/local/lib/python3.10/dist-packages/torch/nn/functional.py:1345: UserWarning:
  warnings.warn(warn_msg)
Epoch: 1 Loss: 0.2929983610663845 Val_Loss: 0.14282329050162185
Epoch: 2 Loss: 0.11167306516834545 Val_Loss: 0.10023150670196206
Epoch: 3 Loss: 0.08669009416963834 Val_Loss: 0.08876393548660906
Epoch: 4 Loss: 0.07547891725945295 Val_Loss: 0.08737611839261517
Epoch: 5 Loss: 0.06766004898735753 Val_Loss: 0.09118692035006895
Epoch: 6 Loss: 0.062466110574673946 Val_Loss: 0.08722920042319367
Epoch: 7 Loss: 0.06020658712990241 Val_Loss: 0.09637939314217543
Epoch: 8 Loss: 0.05766215919957994 Val_Loss: 0.08769152700292669
Epoch: 9 Loss: 0.05625135024997346 Val_Loss: 0.07799628196370527
Epoch: 10 Loss: 0.056682715699112735 Val_Loss: 0.1033533448469157

```

Строим график.

```
# строим график
fig = plt.figure(figsize=(5,5))
plt.plot(np.arange(1,11), loss, "*-",label="Loss")
plt.plot(np.arange(1,11), val_loss,"o-",label="Val Loss")
plt.xlabel("Num of epochs")
plt.legend()
plt.show()
```



Создаем функции атак.

```
# создаем функции атак
def fgsm_attack(input,epsilon,data_grad):
    pert_out = input + epsilon*data_grad.sign()
    pert_out = torch.clamp(pert_out, 0, 1)
    return pert_out

def ifgsm_attack(input,epsilon,data_grad):
    pert_out = input + epsilon*data_grad.sign()
    pert_out = torch.clamp(pert_out, 0, 1)
    return pert_out

def mifgsm_attack(input,epsilon,data_grad):
    iter=10
    decay_factor=1.0
    pert_out = input
    alpha = epsilon/iter
    g=0
    for i in range(iter-1):
        g = decay_factor*g + data_grad/torch.norm(data_grad,p=1)
        pert_out = pert_out + alpha*torch.sign(g)
        pert_out = torch.clamp(pert_out, 0, 1)
        if torch.norm((pert_out-input),p=float('inf')) > epsilon:
            break
    return pert_out
```

Создаем функцию для проверки.


```
# создаем функции проверки
def test(model, device, test_loader, epsilon, attack):
    correct = 0
    adv_examples = []
    for data, target in test_loader:
        data, target = data.to(device), target.to(device)
        data.requires_grad = True
        output = model(data)
        init_pred = output.max(1, keepdim=True)[1]
        if init_pred.item() != target.item():
            continue
        loss = F.nll_loss(output, target)
        model.zero_grad()
        loss.backward()
        data_grad = data.grad.data
        if attack == "fgsm":
            perturbed_data = fgsm_attack(data, epsilon, data_grad)
        elif attack == "ifgsm":
            perturbed_data = ifgsm_attack(data, epsilon, data_grad)
        elif attack == "mifgsm":
            perturbed_data = mifgsm_attack(data, epsilon, data_grad)
        output = model(perturbed_data)
        final_pred = output.max(1, keepdim=True)[1]
        if final_pred.item() == target.item():
            correct += 1
        if (epsilon == 0) and (len(adv_examples) < 5):
            adv_ex = perturbed_data.squeeze().detach().cpu().numpy()
            adv_examples.append( (init_pred.item(), final_pred.item(), adv_ex) )
        else:
            if len(adv_examples) < 5:
                adv_ex = perturbed_data.squeeze().detach().cpu().numpy()
                adv_examples.append( (init_pred.item(), final_pred.item(), adv_ex) )
    final_acc = correct/float(len(test_loader))
    print("Epsilon: {} \t Test Accuracy = {} / {} = {}".format(epsilon, correct, len(test_loader), final_acc))
    return final_acc, adv_examples
```


Выполним построение графиков успешности атак и примеров выполняемых атак в зависимости от степени возмущения epsilon.

```
# строим графики успешности атак и примеры выполненных атак в зависимости от степени возмущения epsilon
epsilons = [0, 0.007, 0.01, 0.02, 0.03, 0.05, 0.1, 0.2, 0.3]
for attack in ("fgsm", "ifgsm", "mifgsm"):
    accuracies = []
    examples = []
    for eps in epsilons:
        acc, ex = test(model, device, test_loader, eps, attack)
        accuracies.append(acc)
        examples.append(ex)
    plt.figure(figsize=(5,5))
    plt.plot(epsilons, accuracies, "-")
    plt.title(attack)
    plt.xlabel("Epsilon")
    plt.ylabel("Accuracy")
    plt.show()

    cnt = 0
    plt.figure(figsize=(8,10))
    for i in range(len(epsilons)):
        for j in range(len(examples[i])):
            cnt += 1
            plt.subplot(len(epsilons), len(examples[0]), cnt)
            plt.xticks([], [])
            plt.yticks([], [])
            if j == 0:
                plt.ylabel("Eps: {}".format(epsilons[i]), fontsize=14)
            orig, adv, ex = examples[i][j]
            plt.title("{} -> {}".format(orig, adv))
            plt.imshow(ex, cmap="gray")
    plt.tight_layout()
    plt.show()
```


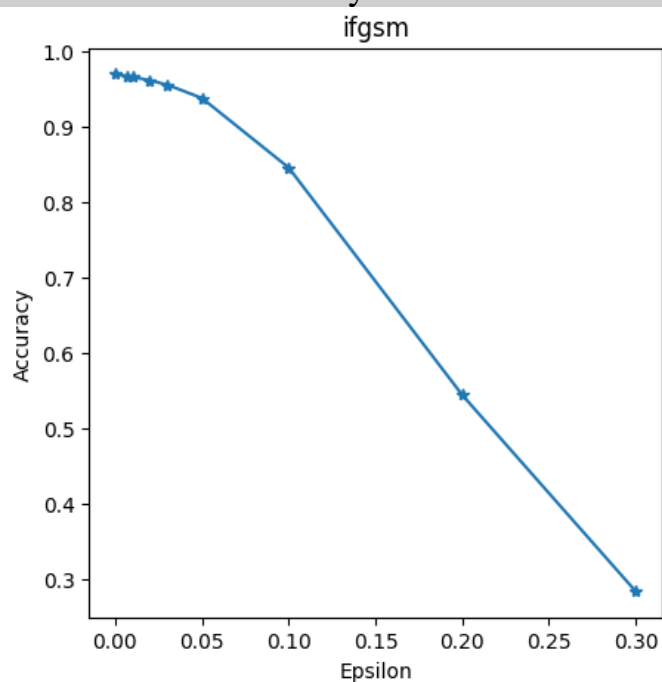
Epsilon: 0	Test Accuracy = 9709 / 10000 = 0.9709
Epsilon: 0.007	Test Accuracy = 9675 / 10000 = 0.9675
Epsilon: 0.01	Test Accuracy = 9654 / 10000 = 0.9654
Epsilon: 0.02	Test Accuracy = 9617 / 10000 = 0.9617
Epsilon: 0.03	Test Accuracy = 9548 / 10000 = 0.9548
Epsilon: 0.05	Test Accuracy = 9364 / 10000 = 0.9364
Epsilon: 0.1	Test Accuracy = 8487 / 10000 = 0.8487
Epsilon: 0.2	Test Accuracy = 5409 / 10000 = 0.5409
Epsilon: 0.3	Test Accuracy = 2789 / 10000 = 0.2789



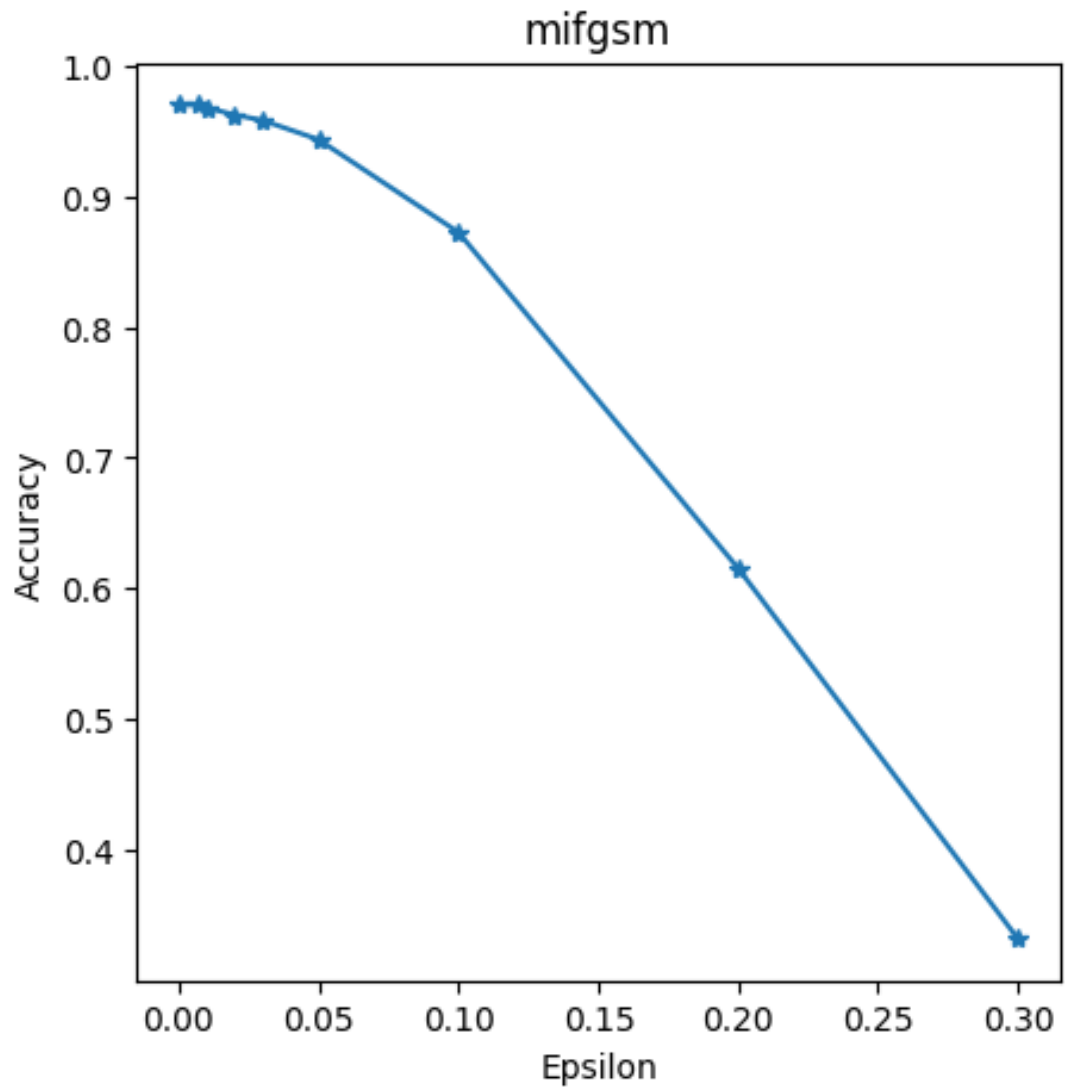



































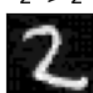

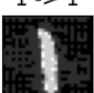









Epsilon: 0	Test Accuracy = 9708 / 10000 = 0.9708
Epsilon: 0.007	Test Accuracy = 9677 / 10000 = 0.9677
Epsilon: 0.01	Test Accuracy = 9666 / 10000 = 0.9666
Epsilon: 0.02	Test Accuracy = 9622 / 10000 = 0.9622
Epsilon: 0.03	Test Accuracy = 9563 / 10000 = 0.9563
Epsilon: 0.05	Test Accuracy = 9387 / 10000 = 0.9387
Epsilon: 0.1	Test Accuracy = 8463 / 10000 = 0.8463

Epsilon: 0.2 Test Accuracy = 5445 / 10000 = 0.5445
Epsilon: 0.3 Test Accuracy = 2837 / 10000 = 0.2837



Epsilon: 0	Test Accuracy = 9700 / 10000 = 0.97
Epsilon: 0.007	Test Accuracy = 9696 / 10000 = 0.9696
Epsilon: 0.01	Test Accuracy = 9676 / 10000 = 0.9676
Epsilon: 0.02	Test Accuracy = 9621 / 10000 = 0.9621
Epsilon: 0.03	Test Accuracy = 9579 / 10000 = 0.9579
Epsilon: 0.05	Test Accuracy = 9432 / 10000 = 0.9432
Epsilon: 0.1	Test Accuracy = 8723 / 10000 = 0.8723
Epsilon: 0.2	Test Accuracy = 6155 / 10000 = 0.6155
Epsilon: 0.3	Test Accuracy = 3324 / 10000 = 0.3324



Eps: 0	3 -> 3 	6 -> 6 	3 -> 3 	5 -> 5 	1 -> 1 
Eps: 0.007	8 -> 8 	7 -> 7 	4 -> 4 	3 -> 3 	0 -> 0 
Eps: 0.01	6 -> 6 	4 -> 4 	8 -> 8 	1 -> 1 	1 -> 1 
Eps: 0.02	4 -> 4 	9 -> 9 	2 -> 2 	5 -> 5 	2 -> 2 
Eps: 0.03	2 -> 2 	4 -> 4 	5 -> 5 	3 -> 3 	5 -> 5 
Eps: 0.05	1 -> 1 	5 -> 5 	5 -> 5 	1 -> 1 	7 -> 7 
Eps: 0.1	5 -> 5 	4 -> 4 	2 -> 2 	2 -> 2 	1 -> 1 
Eps: 0.2	1 -> 1 	4 -> 9 	6 -> 6 	0 -> 0 	2 -> 2 
Eps: 0.3	1 -> 1 	4 -> 4 	2 -> 2 	3 -> 4 	3 -> 0 

Создадим два класса нейросети.

```

# создаем два класса HC
class NetF(nn.Module):
    def __init__(self):
        super(NetF, self).__init__()
        self.conv1 = nn.Conv2d(1, 32, 3, 1)
        self.conv2 = nn.Conv2d(32, 64, 3, 1)
        self.dropout1 = nn.Dropout2d(0.25)
        self.dropout2 = nn.Dropout2d(0.5)
        self.fc1 = nn.Linear(9216, 128)
        self.fc2 = nn.Linear(128, 10)

    def forward(self, x):
        x = self.conv1(x)
        x = F.relu(x)
        x = self.conv2(x)
        x = F.relu(x)
        x = F.max_pool2d(x, 2)
        x = self.dropout1(x)
        x = torch.flatten(x, 1)
        x = self.fc1(x)
        x = F.relu(x)
        x = self.dropout2(x)
        x = self.fc2(x)
        return x

class NetF1(nn.Module):
    def __init__(self):
        super(NetF1, self).__init__()
        self.conv1 = nn.Conv2d(1, 16, 3, 1)
        self.conv2 = nn.Conv2d(16, 32, 3, 1)
        self.dropout1 = nn.Dropout2d(0.25)
        self.dropout2 = nn.Dropout2d(0.5)
        self.fc1 = nn.Linear(4608, 64)
        self.fc2 = nn.Linear(64, 10)

    def forward(self, x):
        x = self.conv1(x)
        x = F.relu(x)
        x = self.conv2(x)
        x = F.relu(x)
        x = F.max_pool2d(x, 2)
        x = self.dropout1(x)
        x = torch.flatten(x, 1)
        x = self.fc1(x)
        x = F.relu(x)
        x = self.dropout2(x)
        x = self.fc2(x)
        return x

```

Создадим функции обучения и тестирования.

```
def fit(model,device,optimizer,scheduler,criterion,train_loader,val_loader,Temp,epochs):
    data_loader = {'train':train_loader,'val':val_loader}
    print("Fitting the model...")
    train_loss,val_loss=[],[]
    for epoch in range(epochs):
        loss_per_epoch,val_loss_per_epoch=0,0
        for phase in ('train','val'):
            for i,data in enumerate(data_loader[phase]):
                input,label = data[0].to(device),data[1].to(device)
                output = model(input)
                output = F.log_softmax(output/Temp,dim=1)
                #calculating loss on the output
                loss = criterion(output,label)
                if phase == 'train':
                    optimizer.zero_grad()
                    #grad calc w.r.t loss func
                    loss.backward()
                    #update weights
                    optimizer.step()
                    loss_per_epoch+=loss.item()
                else:
                    val_loss_per_epoch+=loss.item()
            scheduler.step(val_loss_per_epoch/len(val_loader))
        print("Epoch: {} Loss: {} Val_Loss: {}".format(epoch+1,loss_per_epoch/len(train_loader),val_loss_per_epoch/len(val_loader)))
        train_loss.append(loss_per_epoch/len(train_loader))
        val_loss.append(val_loss_per_epoch/len(val_loader))
    return train_loss,val_loss
```

```
def test(model,device,test_loader,epsilon,Temp,attack):
    correct=0
    adv_examples = []
    for data, target in test_loader:
        data, target = data.to(device), target.to(device)
        data.requires_grad = True
        output = model(data)
        output = F.log_softmax(output/Temp,dim=1)
        init_pred = output.max(1, keepdim=True)[1]
        if init_pred.item() != target.item():
            continue
        loss = F.nll_loss(output, target)
        model.zero_grad()
        loss.backward()
        data_grad = data.grad.data
        if attack == "fgsm":
            perturbed_data = fgsm_attack(data,epsilon,data_grad)
        elif attack == "ifgsm":
            perturbed_data = ifgsm_attack(data,epsilon,data_grad)
        elif attack == "mifgsm":
            perturbed_data = mifgsm_attack(data,epsilon,data_grad)
        output = model(perturbed_data)
        final_pred = output.max(1, keepdim=True)[1]
        if final_pred.item() == target.item():
            correct += 1
            if (epsilon == 0) and (len(adv_examples) < 5):
                adv_ex = perturbed_data.squeeze().detach().cpu().numpy()
                adv_examples.append( (init_pred.item(), final_pred.item(), adv_ex) )
            else:
                if len(adv_examples) < 5:
                    adv_ex = perturbed_data.squeeze().detach().cpu().numpy()
                    adv_examples.append( (init_pred.item(), final_pred.item(), adv_ex) )
        final_acc = correct/float(len(test_loader))
    print("Epsilon: {} \t Test Accuracy = {} / {} = {}".format(epsilon, correct, len(test_loader), final_acc))
    return final_acc,adv_examples
```

Зададим функцию защиты сетей.

```

def defense(device,train_loader,val_loader,test_loader,epochs,Temp,epsilons):
    modelF = NetF().to(device)
    optimizerF = optim.Adam(modelF.parameters(),lr=0.0001, betas=(0.9, 0.999))
    schedulerF = optim.lr_scheduler.ReduceLROnPlateau(optimizerF, mode='min', factor=0.1, patience=3)
    modelF1 = NetF1().to(device)
    optimizerF1 = optim.Adam(modelF1.parameters(),lr=0.0001, betas=(0.9, 0.999))
    schedulerF1 = optim.lr_scheduler.ReduceLROnPlateau(optimizerF1, mode='min', factor=0.1, patience=3)
    criterion = nn.NLLLoss()
    lossF,val_lossF=fit(modelF,device,optimizerF,schedulerF,criterion,train_loader,val_loader,Temp,epochs)
    fig = plt.figure(figsize=(5,5))
    plt.plot(np.arange(1,epochs+1), lossF, "--",label="Loss")
    plt.plot(np.arange(1,epochs+1), val_lossF,"o-",label="Val Loss")
    plt.title("Network F")
    plt.xlabel("Num of epochs")
    plt.legend()
    plt.show()
    #converting target labels to soft labels
    for data in train_loader:
        input, label = data[0].to(device),data[1].to(device)
        softlabel = F.log_softmax(modelF(input),dim=1)
        data[1] = softlabel
    lossF1,val_lossF1=fit(modelF1,device,optimizerF1,schedulerF1,criterion,train_loader,val_loader,Temp,epochs)
    fig = plt.figure(figsize=(5,5))
    plt.plot(np.arange(1,epochs+1), lossF1, "--",label="Loss")
    plt.plot(np.arange(1,epochs+1), val_lossF1,"o-",label="Val Loss")
    plt.title("Network F'")
    plt.xlabel("Num of epochs")
    plt.legend()
    plt.show()
    model = NetF1().to(device)
    model.load_state_dict(modelF1.state_dict())
    for attack in ("fgsm","ifgsm","mifgsm"):
        accuracies = []
        examples = []
        for eps in epsilons:
            acc, ex = test(model,device,test_loader,eps,"fgsm")
            accuracies.append(acc)
            examples.append(ex)
        plt.figure(figsize=(5,5))
        plt.plot(epsilons, accuracies, "--")
        plt.title(attack)
        plt.xlabel("Epsilon")
        plt.ylabel("Accuracy")
        plt.show()

```

```

cnt = 0
plt.figure(figsize=(8,10))
for i in range(len(epsilons)):
    for j in range(len(examples[i])):
        cnt += 1
        plt.subplot(len(epsilons),len(examples[0]),cnt)
        plt.xticks([], [])
        plt.yticks([], [])
        if j == 0:
            plt.ylabel("Eps: {}".format(epsilons[i]), fontsize=14)
        orig,adv,ex = examples[i][j]
        plt.title("{} -> {}".format(orig, adv))
        plt.imshow(ex, cmap="gray")
plt.tight_layout()
plt.show()

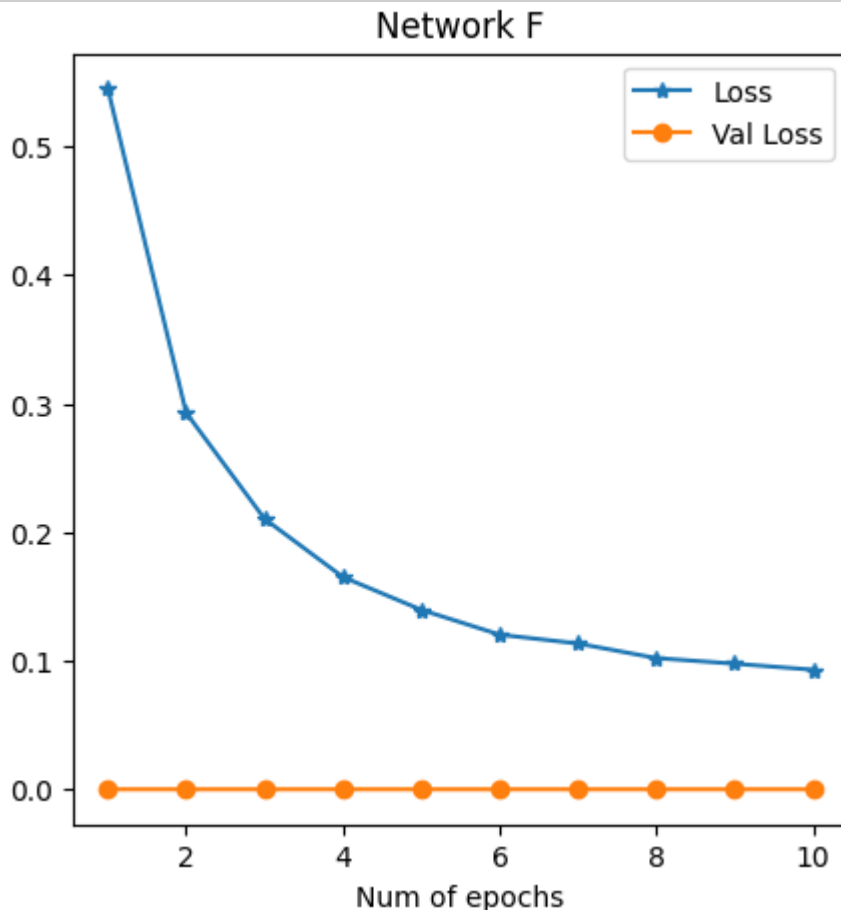
```

Получим результаты оценки защищенных сетей.

```
# получаем результаты оценки защищенных сетей
Temp=100
epochs=10
epsilons=[0,0.007,0.01,0.02,0.03,0.05,0.1,0.2,0.3]
defense(device,train_loader,val_loader,test_loader,epochs,Temp,epsilons)
```

Fitting the model...


Epoch: 1 Loss: 0.5448293192800642 Val_Loss: 7.73291103541851e-06
Epoch: 2 Loss: 0.2927316698674613 Val_Loss: 1.0102411732077599e-06
Epoch: 3 Loss: 0.2102989830346451 Val_Loss: 6.601312197744847e-07
Epoch: 4 Loss: 0.16505023587049214 Val_Loss: 7.533462459105067e-06
Epoch: 5 Loss: 0.13952160401099645 Val_Loss: 4.4946352019906044e-06
Epoch: 6 Loss: 0.12023352776694962 Val_Loss: 1.9566755509003998e-07
Epoch: 7 Loss: 0.11333786990761777 Val_Loss: 1.9792682214756497e-06
Epoch: 8 Loss: 0.1019939355488022 Val_Loss: 4.203413613140583e-07
Epoch: 9 Loss: 0.0976399402378025 Val_Loss: 2.3648618487641217e-06
Epoch: 10 Loss: 0.09295012344814668 Val_Loss: 7.459939115506131e-08



Fitting the model...


Epoch: 1 Loss: 0.6363280130847226 Val_Loss: 4.991228529252112e-06
Epoch: 2 Loss: 0.42811888280111854 Val_Loss: 6.821165650617332e-05
Epoch: 3 Loss: 0.3726626193641786 Val_Loss: 0.0001593279018998146
Epoch: 4 Loss: 0.3215511648770174 Val_Loss: 9.6123790089041e-05
Epoch: 5 Loss: 0.2729614160263343 Val_Loss: 1.9890321791172028e-05


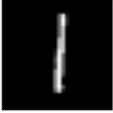











































Epoch: 6 Loss: 0.2336730380693122 Val_Loss: 1.418686844408512e-05
Epoch: 7 Loss: 0.22432161142330115 Val_Loss: 5.9460494473751167e-05
Epoch: 8 Loss: 0.22250396553932653 Val_Loss: 2.616085112094879e-07
Epoch: 9 Loss: 0.21347470993006776 Val_Loss: 2.3429323930758982e-05
Epoch: 10 Loss: 0.21204543162817313 Val_Loss: 1.9299010455142708e-08



Epsilon: 0	Test Accuracy = 9114 / 10000 = 0.9114
Epsilon: 0.007	Test Accuracy = 9015 / 10000 = 0.9015
Epsilon: 0.01	Test Accuracy = 9007 / 10000 = 0.9007
Epsilon: 0.02	Test Accuracy = 8973 / 10000 = 0.8973
Epsilon: 0.03	Test Accuracy = 8935 / 10000 = 0.8935
Epsilon: 0.05	Test Accuracy = 8718 / 10000 = 0.8718
Epsilon: 0.1	Test Accuracy = 8046 / 10000 = 0.8046
Epsilon: 0.2	Test Accuracy = 4897 / 10000 = 0.4897
Epsilon: 0.3	Test Accuracy = 2008 / 10000 = 0.2008
Epsilon: 0	Test Accuracy = 9128 / 10000 = 0.9128
Epsilon: 0.007	Test Accuracy = 9018 / 10000 = 0.9018
Epsilon: 0.01	Test Accuracy = 9038 / 10000 = 0.9038
Epsilon: 0.02	Test Accuracy = 8954 / 10000 = 0.8954
Epsilon: 0.03	Test Accuracy = 8897 / 10000 = 0.8897
Epsilon: 0.05	Test Accuracy = 8735 / 10000 = 0.8735
Epsilon: 0.1	Test Accuracy = 8077 / 10000 = 0.8077
Epsilon: 0.2	Test Accuracy = 4935 / 10000 = 0.4935

Epsilon: 0.3	Test Accuracy = 1992 / 10000 = 0.1992
Epsilon: 0	Test Accuracy = 9102 / 10000 = 0.9102
Epsilon: 0.007	Test Accuracy = 9030 / 10000 = 0.903
Epsilon: 0.01	Test Accuracy = 9034 / 10000 = 0.9034
Epsilon: 0.02	Test Accuracy = 8948 / 10000 = 0.8948
Epsilon: 0.03	Test Accuracy = 8951 / 10000 = 0.8951
Epsilon: 0.05	Test Accuracy = 8733 / 10000 = 0.8733
Epsilon: 0.1	Test Accuracy = 8047 / 10000 = 0.8047
Epsilon: 0.2	Test Accuracy = 4898 / 10000 = 0.4898
Epsilon: 0.3	Test Accuracy = 1969 / 10000 = 0.1969



Eps: 0	1 -> 1 	1 -> 1 	0 -> 0 	1 -> 1 	2 -> 2 
Eps: 0.007	6 -> 6 	8 -> 8 	1 -> 1 	7 -> 4 	2 -> 2 
Eps: 0.01	1 -> 1 	6 -> 6 	9 -> 9 	4 -> 4 	3 -> 3 
Eps: 0.02	8 -> 8 	3 -> 8 	7 -> 8 	2 -> 2 	4 -> 4 
Eps: 0.03	9 -> 9 	0 -> 0 	0 -> 0 	8 -> 8 	3 -> 8 
Eps: 0.05	5 -> 5 	5 -> 5 	7 -> 7 	3 -> 3 	3 -> 3 
Eps: 0.1	3 -> 3 	6 -> 6 	3 -> 3 	9 -> 9 	5 -> 5 
Eps: 0.2	9 -> 9 	5 -> 0 	7 -> 5 	9 -> 3 	2 -> 2 
Eps: 0.3	3 -> 5 	5 -> 5 	1 -> 2 	4 -> 1 	7 -> 4 

Вывод

Без использования защитных механизмов (при $\text{eps}=0$) точность модели высока, и она составляет 91.28%.

С увеличением значения eps точность модели ухудшается, особенно при больших значениях. Это свидетельствует о том, что модель чувствительна к атакам с добавлением шума (например, FGSM, iFGSM, MI-FGSM).

После применения механизма защиты (путем обучения модели на защищенных данных) наблюдается улучшение в стойкости модели к атакам.

Точность модели при различных значениях ϵ ps становится выше по сравнению с беззащитным вариантом. Особенно заметно улучшение на больших значениях ϵ ps. Однако, даже с использованием защиты, модель не обладает абсолютной стойкостью к атакам, и точность все равно снижается с увеличением ϵ ps.

Графики показывают, как меняется точность модели в зависимости от степени атаки (ϵ ps) для различных методов (FGSM, iFGSM, MI-FGSM). С увеличением ϵ ps точность модели снижается, что является ожидаемым результатом при увеличении степени атаки.

Визуализация примеров атак позволяет увидеть, как изменяется внешний вид изображений после атаки для различных значений ϵ ps.

Можно заметить, что с увеличением ϵ ps изображения становятся все более искаженными и трудными для классификации моделью.